

Highlights

Complexity-Constrained Few-Shot Digital Twin Model Instantiation for Heterogeneous Multi-Center Computing Networks

- RCF-DTI instantiates digital twin models from limited target data.
- A compact candidate set supports several source-trained architectures.
- Models above the assigned complexity limits are removed before adaptation.
- A fixed reference reduces changes caused by small validation gains.
- RCF-DTI lowers mean squared error while every selected model meets both limits.

Complexity-Constrained Few-Shot Digital Twin Model Instantiation for Heterogeneous Multi-Center Computing Networks

ARTICLE INFO

Keywords:

digital twin
few-shot learning
model instantiation
architecture selection
model complexity
edge-cloud computing

ABSTRACT

Shared edge-cloud digital twin platforms maintain target-specific models for heterogeneous computing centers. Because the platform hosts several model instances, it assigns each instance limits on inference computation and model size. A new or reconfigured center may provide only a few observations. The platform must therefore select a suitable architecture and adapt its parameters from limited data while meeting both limits. Existing methods usually adapt one fixed architecture or compare several architectures without jointly handling model limits and unreliable model selection from a small validation set. We propose the Reference-Regularized Complexity-Constrained Few-Shot Digital Twin Instantiation algorithm (RCF-DTI). It removes candidate models that exceed the assigned limits, adapts the remaining source-trained candidates with the same update budget, and selects among them using a fixed reference candidate. An alternative is selected only when it gives sufficient validation improvement. Experiments show that RCF-DTI reduces mean squared error by 14.60% relative to pretraining and fine-tuning, while every selected model satisfies both assigned limits.

1. Introduction

Digital twin platforms maintain virtual model instances of physical systems to support monitoring, prediction, analysis, and decision making [1–4]. In a heterogeneous multi-center computing network, a shared edge-cloud platform maintains a target-specific digital twin model instance for each physical center. The centers may differ in monitored variables, workload dynamics, and data distributions. They may therefore require different model architectures. Since the platform serves several model instances at the same time, it assigns each instance limits on inference computation and model size according to its hosting conditions and service requirements. We measure these two properties by estimated operation count and parameter count, respectively.

A new or reconfigured center may provide only a few observations. An existing platform instance may not fit the target center or satisfy the complexity limits assigned to the target instance. Direct reuse is therefore not always suitable. The platform must select a suitable architecture and adapt its parameters under both limits. We call this joint process target-specific digital twin model instantiation. The main challenge is to make this decision reliably from limited target data. Figure 1 illustrates this setting.




Figure 1: Background of platform-side target-specific digital twin model instantiation in a heterogeneous multi-center computing network. The shared edge–cloud platform maintains model instances for heterogeneous physical centers. For a new or reconfigured target center, only a few observations may be available, and an existing instance may not fit the target center or satisfy the assigned limits on estimated operation count and parameter count. The platform must therefore determine a target-specific digital twin instance.

Existing studies address the main parts of this problem separately. Digital twin research studies system architecture, model construction, lifecycle management, and service control, but it usually assumes a predefined model structure or one general model for a class of entities [1, 3, 5, 7]. Few-shot and meta-learning methods transfer knowledge to new tasks with limited data, but most of them retain one fixed architecture [9–11, 27]. Deployment-aware architecture search considers model accuracy and execution cost [12, 13]. Few-shot, zero-shot, training-free, and multi-objective methods reduce repeated training or improve architecture ranking [14–16, 28]. However, they do not focus on reliable target-specific selection from a small validation set. Existing methods therefore provide limited support for jointly selecting an architecture and adapting its parameters under platform-assigned complexity limits. They also do not directly control changes caused by small validation differences.

To address this problem, we propose the Reference-Regularized Complexity-Constrained Few-Shot Digital Twin Instantiation algorithm, termed RCF-DTI. It formulates target-specific instantiation as a constrained two-level optimization problem. The inner level adapts each feasible candidate. The outer level selects one candidate using a reference-regularized validation score. An offline stage constructs a compact set of source-trained candidates with different architectures. For each target center, RCF-DTI removes infeasible candidates, adapts the remaining candidates with the same update budget, and selects a non-reference candidate only when it gives sufficient validation improvement. The output is the selected architecture and its adapted parameters. Multi-center workload prediction [17–19] is used only for evaluation.

The main contributions are as follows.

- We formulate target-specific digital twin model instantiation as a constrained two-level optimization problem. The formulation jointly considers few-shot parameter adaptation, architecture selection, and hard limits on estimated operation count and parameter count. A reference-regularized objective requires sufficient validation improvement over a fixed reference candidate.
- We design RCF-DTI to solve the constrained problem over a finite candidate bank. It separates complexity filtering from parameter adaptation because model complexity is fixed by architecture. It also uses a common adaptation budget to make candidate losses comparable and a calibrated reference rule to avoid replacements based on small validation differences.
- We establish output feasibility and the minimum validation improvement of accepted alternatives. Under bounded validation error, we derive a sufficient condition for an accepted alternative to improve the expected loss on future target observations. We also give the target-stage computational complexity. Experiments confirm improved prediction accuracy and full satisfaction of the assigned limits.

The remainder of this paper is organized as follows. Section 2 reviews related work. Section 3 defines the system model and problem. Section 4 presents RCF-DTI. Section 5 reports the experiments, and Section 6 concludes the paper.

2. Related Work

2.1. Digital Twin Modeling and Platform Deployment

Digital twins maintain virtual representations of physical entities through continuous data exchange. They support monitoring, prediction, simulation, and control [1–3]. In industrial and networked systems, data-driven models capture system states and support management decisions [4–6]. Related time-series studies estimate missing observations with probabilistic models [20] and transfer forecasting knowledge across related data sources [21]. Digital twin networks further extend this idea to network elements, computing nodes, and services. Their models may be hosted by cloud or edge platforms while remaining associated with individual physical entities [7, 8]. Service orchestration and computation offloading also consider task placement, delay, and resource use in edge–cloud systems [29, 30].

These studies establish the roles of data-driven modeling and platform-side deployment. However, they mainly focus on system architecture, model construction, lifecycle management, data processing, or service control. They usually use a predefined model structure or one general model for a class of entities. They therefore provide limited support for selecting different architectures for newly observed heterogeneous centers under instance-specific complexity limits.

2.2. Few-Shot Learning and Meta-Learning

Few-shot learning aims to learn a target task from a small number of target examples. Meta-learning further learns transferable knowledge, initializations, or update rules from related tasks [9, 10]. These methods have been applied to digital twin construction, context adaptation, and cross-domain transfer [11, 22, 23]. Prompt-based and memory-based methods improve transfer from limited target data [24, 25]. Online and transfer-learning studies further address sequential tasks and workload changes [26, 27].

These methods improve parameter adaptation from limited target data. However, most of them retain one fixed architecture during target adaptation. They do not jointly determine the architecture and parameters of a target-specific digital twin model. They also do not enforce platform-assigned model complexity limits during this joint decision.

2.3. Deployment-Aware Architecture Selection

Neural architecture search selects a model structure from a predefined architecture space. Platform-aware and hardware-aware methods consider model accuracy and execution cost during search [12, 13]. Few-shot neural architecture search reduces repeated candidate training by sharing knowledge among architectures [14]. Zero-shot and training-free methods estimate architecture quality without fully training each candidate [15, 16]. Multi-objective search further considers several performance and complexity measures [28].

These methods reduce architecture search cost, but they mainly focus on search efficiency or architecture ranking. They are not designed for target-specific digital twin instantiation with both a small support set and a small validation set. They also do not jointly handle hard instance-specific limits and unreliable choices caused by small validation differences.

Overall, existing work provides limited support for jointly handling architecture selection, few-shot parameter adaptation, hard model complexity limits, and unreliable comparison based on limited validation data. RCF-DTI addresses this gap by removing models that exceed the limits, adapting the remaining models with the same update budget, and requiring sufficient validation improvement before selecting an alternative.

3. System Model and Problem Formulation

3.1. Platform Setting

We consider a shared edge–cloud digital twin platform that maintains model instances for heterogeneous computing centers. Let C_{src} denote the source centers that provide historical data for offline training. During each instantiation process, we consider one new or reconfigured target center c with limited observations.

Each physical center is associated with one model instance hosted by the shared platform. The centers may differ in monitored variables, workload patterns, and data distributions. They may therefore require different model architectures.

Let $\mathcal{A}$ denote the predefined architecture space. For an architecture $a \in \mathcal{A}$, let $f_{a,\theta}$ denote a data-driven digital twin model with parameter vector θ . The digital twin model instance for target center c is represented by

$$\text{DT}_c = (a_c, \theta_c), \quad (1)$$

where a_c is the selected model architecture and θ_c is its target-adapted parameter vector. Model instantiation must therefore determine both the architecture and the model parameters.

3.2. Candidate Models and Complexity Limits

The target observations are divided in chronological order into a support set $\mathcal{D}_c^{\text{sup}}$ and a validation set $\mathcal{D}_c^{\text{val}}$. The support set is used for target adaptation. The validation set is used for candidate selection.

A candidate model is denoted by $q = (a_q, \theta_q^0)$. Here, a_q is the candidate architecture and θ_q^0 is its source-trained initialization. We call a finite set of candidate models a candidate bank. Let $\mathcal{Q}$ denote the candidate bank constructed before target instantiation. Different candidates may use different architectures or different source-trained initializations.

The platform assigns two complexity limits to the digital twin instance of target center c . The first limit controls the computation required for one inference. The second limit controls the model size. For architecture a_q , let $F_c(a_q)$ denote its estimated operation count under the target input setting. Let $P(a_q)$ denote its parameter count. Their platform-assigned upper bounds are denoted by B_c^F and B_c^P , respectively.

The feasible candidate set is

$$\mathcal{Q}_c^{\text{fea}} = \{q \in \mathcal{Q} \mid F_c(a_q) \leq B_c^F, P(a_q) \leq B_c^P\}. \quad (2)$$

A candidate belongs to $\mathcal{Q}_c^{\text{fea}}$ only when it satisfies both limits. The two limits are hard constraints. They are not added to the prediction loss as penalty terms.

One candidate q^{ref} is fixed as the reference candidate before target instantiation. We assume that $q^{\text{ref}} \in \mathcal{Q}_c^{\text{fea}}$ for each target center to which the algorithm is applied. The reference candidate provides a fixed default choice when an alternative candidate does not provide enough validation improvement.

3.3. Problem Formulation

Let $\mathcal{L}(\mathcal{D}; a, \theta)$ denote the mean task loss of model $f_{a,\theta}$ on dataset $\mathcal{D}$. The task loss used in the experiments is given in Section 5.

Let T denote the fixed number of target parameter updates. For each feasible candidate q , let $\theta_{c,q}^{(T)}$ denote its parameter vector after T updates on $\mathcal{D}_c^{\text{sup}}$. The concrete update rule is given in Section 4. The post-adaptation validation loss is

$$J_c(q) = \mathcal{L}(\mathcal{D}_c^{\text{val}}; a_q, \theta_{c,q}^{(T)}). \quad (3)$$

Directly returning the candidate with the lowest value of $J_c(q)$ may be unreliable when the validation set is small. A small difference on the validation set may disappear on unseen target data. We therefore add a relative selection cost when a non-reference candidate is selected. We call this cost reference regularization.

Let $\tau \in [0, 1)$ denote the required relative improvement over the reference candidate. The regularized selection score is defined as

$$S_c(q; \tau) = \begin{cases} J_c(q), & q = q^{\text{ref}}, \\ J_c(q) + \tau J_c(q^{\text{ref}}), & q \neq q^{\text{ref}}. \end{cases} \quad (4)$$

The added term is used only for candidate selection. It does not change the parameter adaptation loss. A larger value of τ requires a larger validation improvement before the algorithm selects a non-reference candidate.

The target-specific instantiation problem is

$$q_c^* \in \arg \min_{q \in \mathcal{Q}_c^{\text{fea}}} S_c(q; \tau). \quad (5)$$

The problem has two levels. For each feasible candidate, the inner level adapts its continuous parameter vector using the target support set. The outer level selects one candidate under the complexity limits using the score in Eq. (4). We therefore formulate target-specific model instantiation as a constrained two-level optimization problem.

4. RCF-DTI Method

4.1. Overview

RCF-DTI uses a candidate bank prepared offline together with a fixed reference candidate and a calibrated selection threshold. For each target center, the algorithm receives the target data and assigned complexity limits. It first removes candidates that violate either limit, then adapts the feasible candidates with a fixed update budget, and finally performs reference-regularized selection. The selected architecture and its adapted parameters form the target-specific digital twin instance. Figure 2 summarizes the overall RCF-DTI workflow and its relation to digital twin operation. The offline bank preparation and the target-side stages are detailed in the following subsections.




Figure 2: Overview of RCF-DTI. A candidate bank prepared offline and the target data with assigned limits are provided to the algorithm. RCF-DTI performs complexity filtering, fixed-budget adaptation, and reference-regularized selection. The selected candidate forms the target-specific digital twin instance, which exchanges monitoring data and state estimates with the target center during digital twin operation.

4.2. Offline Bank Construction

The candidate-bank input in Figure 2 is prepared offline in two steps. First, a compact architecture set is retained from the predefined space using separate development data.

Let $\mathcal{A}_{\text{bank}} = \{a_m\}_{m=1}^M$ denote the retained architecture set, where M is the number of retained architectures. Second, let $\mathcal{D}_{\text{src}}$ denote the pooled source-center training data. For each retained architecture a_m , the offline stage trains one or more source initializations. The r -th initialization is trained by

$$\theta_{m,r}^0 \in \arg \min_{\theta} \mathcal{L}(\mathcal{D}_{\text{src}}; a_m, \theta), \quad m = 1, \dots, M, \quad r = 1, \dots, R_m, \quad (6)$$

where θ is the parameter variable during source training and R_m is the number of stored initializations for architecture a_m . In practice, the source optimizer gives an approximate solution to Eq. (6).

Each pair $q_{m,r} = (a_m, \theta_{m,r}^0)$ forms one candidate. The complete bank is $\mathcal{Q} = \{q_{m,r}\}$ for $m = 1, \dots, M$ and $r = 1, \dots, R_m$. Because each initialization is trained with its own architecture, it is called an architecture-matched source initialization.

The reference candidate q^{ref} is fixed from this bank before target instantiation and is not changed across target centers. Figure 3 details the offline construction of the candidate bank shown in Figure 2.




Figure 3: Offline construction of the candidate bank used by RCF-DTI. A compact architecture set $\mathcal{A}_{\text{bank}}$ is retained from the predefined architecture space using separate development data. For each retained architecture a_m , pooled source-center data are used to train one or more architecture-matched source initializations $\theta^0_{m,r}$. Each pair $q_{m,r} = (a_m, \theta^0_{m,r})$ forms one candidate, and one bank member is fixed as the reference candidate q^{ref} before target instantiation.

4.3. Candidate Filtering

For target center c , RCF-DTI computes the feasible set using Eq. (2). Filtering compares the estimated operation count $F_c(a_q)$ and parameter count $P(a_q)$ with their assigned upper bounds B_c^F and B_c^P , respectively. It does not use the target support or validation data. Because architecture complexity does not change during parameter adaptation, a candidate that violates either limit is removed before target updates. This avoids unnecessary adaptation. The feasible reference candidate ensures at least one valid output. Figure 4 illustrates representative complexity-filtering outcomes.




Figure 4: Complexity filtering under platform-assigned instance limits. A candidate q remains feasible only when its estimated operation count $F_c(a_q)$ and parameter count $P(a_q)$ do not exceed the assigned upper bounds B_c^F and B_c^P , respectively. Violation of either limit makes the candidate infeasible. The fixed reference candidate q^{ref} is highlighted in blue.

4.4. Target Adaptation

Each feasible candidate starts from its architecture-matched source initialization. Let $\eta > 0$ denote the target learning rate, and let $G > 0$ denote the maximum gradient norm. The parameters are updated by

$$\begin{aligned} \theta_{c,q}^{(0)} &= \theta_q^0, \\ \theta_{c,q}^{(t+1)} &= \theta_{c,q}^{(t)} - \eta \text{Clip}_G \left(\nabla_{\theta} \mathcal{L} \left(\mathcal{D}_c^{\text{sup}}; a_q, \theta_{c,q}^{(t)} \right) \right), \\ t &= 0, \dots, T-1. \end{aligned} \quad (7)$$

Here, t is the update index, and ∇_{θ} denotes the gradient with respect to the model parameters. The function $\text{Clip}_G(\cdot)$ limits the Euclidean norm of the gradient to G . It leaves the gradient unchanged when the norm does not exceed G and rescales it otherwise. The architecture a_q remains fixed during all updates.

All feasible candidates use the same support set, task loss, optimizer, learning rate, gradient clipping rule, and update budget T . Candidate-specific early stopping is not used. These common settings make the post-adaptation losses comparable.

After the final update, the algorithm computes $J_c(q)$ using Eq. (3). The validation set is used only for model selection. It is not used for parameter updates.

4.5. Candidate Selection

Let $Q_c^{\text{alt}} = Q_c^{\text{fea}} \setminus \{q^{\text{ref}}\}$ denote the set of feasible alternatives. If this set is empty, the algorithm returns the reference candidate.

Otherwise, let $\tilde{q}_c$ denote a feasible alternative with the lowest post-adaptation validation loss:

$$\tilde{q}_c \in \arg \min_{q \in Q_c^{\text{alt}}} J_c(q).$$

The final candidate is selected by

$$q_c^* = \begin{cases} \tilde{q}_c, & J_c(\tilde{q}_c) \leq (1 - \tau)J_c(q^{\text{ref}}), \\ q^{\text{ref}}, & \text{otherwise.} \end{cases} \quad (8)$$

For the computed post-adaptation losses, Eq. (8) gives a minimizer of the selection problem in Eq. (5). All alternatives receive the same added cost $\tau J_c(q^{\text{ref}})$. Therefore, the alternative with the lowest value of $J_c(q)$ also has the lowest regularized score among the alternatives. It is selected only when the condition in Eq. (8) is satisfied.

Figure 5 illustrates the fixed-budget adaptation and reference-regularized selection stages.




Figure 5: Fixed-budget target adaptation and reference-regularized selection. (a) Each feasible candidate starts from its architecture-matched source initialization and updates only its parameters for T steps using the same support data and update procedure. (b) The best feasible alternative is selected only when $J_c(\tilde{q}_c) \leq (1 - \tau)J_c(q^{\text{ref}})$; otherwise, the fixed reference candidate is retained.

Algorithm 1 summarizes RCF-DTI for one target center.

Algorithm 1 RCF-DTI for Target-Specific Model Instantiation

Require: Candidate bank $\mathcal{Q}$, reference candidate q^{ref} , support set $\mathcal{D}_c^{\text{sup}}$, and validation set $\mathcal{D}_c^{\text{val}}$
Require: Limits B_c^F and B_c^P , update budget T , learning rate η , gradient bound G , and calibrated threshold τ
Ensure: Target-specific digital twin model DT_c^*

- 1: Compute $\mathcal{Q}_c^{\text{fea}}$ using Eq. (2)
- 2: **for all** $q \in \mathcal{Q}_c^{\text{fea}}$ **do**
- 3: $\theta_{c,q}^{(0)} \leftarrow \theta_q^0$
- 4: **for** $t = 0$ to $T - 1$ **do**
- 5: $g \leftarrow \nabla_{\theta} \mathcal{L}(\mathcal{D}_c^{\text{sup}}; a_q, \theta_{c,q}^{(t)})$
- 6: $g \leftarrow \text{Clip}_G(g)$
- 7: $\theta_{c,q}^{(t+1)} \leftarrow \theta_{c,q}^{(t)} - \eta g$
- 8: **end for**
- 9: Compute $J_c(q)$ using Eq. (3)
- 10: **end for**
- 11: **for all** $q \in \mathcal{Q}_c^{\text{fea}}$ **do**
- 12: Compute $S_c(q; \tau)$ using Eq. (4)
- 13: **end for**
- 14: Choose $q_c^* \in \arg \min_{q \in \mathcal{Q}_c^{\text{fea}}} S_c(q; \tau)$
- 15: $\text{DT}_c^* \leftarrow (a_{q_c^*}, \theta_{c,q_c^*}^{(T)})$
- 16: **return** DT_c^*

4.6. Threshold Calibration

The threshold τ is fixed before target instantiation. It is selected from a finite set $\mathcal{T}$ using a separate set of development cases, denoted by $\mathcal{C}_{\text{dev}}$. Each development case contains support, validation, and check data. The support and validation data are used by RCF-DTI. The check data are used only for threshold calibration.

For each case $d \in \mathcal{C}_{\text{dev}}$, let $h_d(\tau) = 1$ when RCF-DTI selects a non-reference candidate whose check loss is higher than that of the reference candidate. Otherwise, let $h_d(\tau) = 0$. The empirical harmful-selection rate is

$$\hat{R}_{\text{harm}}(\tau) = \frac{1}{|\mathcal{C}_{\text{dev}}|} \sum_{d \in \mathcal{C}_{\text{dev}}} h_d(\tau). \quad (9)$$

Calibration also requires positive development gains. Let $\hat{G}_{\text{ref}}(\tau)$ denote the mean relative check-loss reduction over the fixed reference candidate. To distinguish gains from architecture selection from gains caused only by a different source initialization, we define a reference-architecture control. This control applies the same selection rule only to candidates that use the reference architecture. Let $\hat{G}_{\text{arch}}(\tau)$ denote the mean additional reduction over this control.

Let $\underline{G}_{\text{ref}}(\tau)$ and $\underline{G}_{\text{arch}}(\tau)$ denote the lower confidence bounds of the two gains. Let ρ denote the allowed harmful-selection rate, and let γ_{ref} and γ_{arch} denote the required minimum gains. The eligible threshold set and the calibrated threshold are

$$\begin{aligned} \mathcal{T}_{\text{elig}} = \{ \tau \in \mathcal{T} \mid & \hat{R}_{\text{harm}}(\tau) \leq \rho, \\ & \hat{G}_{\text{ref}}(\tau) \geq \gamma_{\text{ref}}, \quad \underline{G}_{\text{ref}}(\tau) > 0, \\ & \hat{G}_{\text{arch}}(\tau) \geq \gamma_{\text{arch}}, \quad \underline{G}_{\text{arch}}(\tau) > 0 \}, \\ \tau^* = \min \mathcal{T}_{\text{elig}}. \end{aligned} \quad (10)$$

We assume that $\mathcal{T}_{\text{elig}}$ is nonempty. The smallest eligible threshold is used because it satisfies the development criteria without requiring a larger validation improvement than necessary. The tested values, gain requirements, and confidence procedure are reported in Section 5.1.

This calibration controls the observed behavior on the development cases. It does not guarantee improvement for every unseen target center.

4.7. Algorithm Analysis and Complexity

Complexity feasibility. Algorithm 1 selects a candidate only from $\mathcal{Q}_c^{\text{fea}}$. Therefore, its output always satisfies

$$F_c(a_{q_c^*}) \leq B_c^F, \quad P(a_{q_c^*}) \leq B_c^P.$$

This property follows directly from the hard filtering step.

Minimum validation improvement. Suppose that the algorithm returns a non-reference candidate and that $J_c(q^{\text{ref}}) > 0$. Equation (8) gives

$$\frac{J_c(q^{\text{ref}}) - J_c(q_c^*)}{J_c(q^{\text{ref}})} \geq \tau.$$

Thus, every accepted alternative reaches at least the required relative validation improvement.

Selection under bounded validation error. Let $R_c(q)$ denote the expected loss of the adapted candidate q on future observations from target center c . Suppose that the validation loss approximates this expected loss with error at most ϵ_c for the selected alternative and the reference candidate:

$$|R_c(q) - J_c(q)| \leq \epsilon_c, \quad q \in \{q_c^*, q^{\text{ref}}\}.$$

If RCF-DTI selects a non-reference candidate, then Eq. (8) gives

$$R_c(q_c^*) - R_c(q^{\text{ref}}) \leq -\tau J_c(q^{\text{ref}}) + 2\epsilon_c. \quad (11)$$

Therefore, when $\tau J_c(q^{\text{ref}}) > 2\epsilon_c$, the selected alternative has a lower expected loss than the reference candidate on future target observations. This is a sufficient condition under the stated validation error bound. It does not provide an unconditional guarantee for every target center.

Target-stage computation. Let C_q^{grad} denote the cost of one gradient update for candidate q , and let C_q^{val} denote its validation cost. The target-stage computation is

$$O\left(|\mathcal{Q}| + \sum_{q \in \mathcal{Q}_c^{\text{fea}}} \left(T C_q^{\text{grad}} + C_q^{\text{val}}\right)\right). \quad (12)$$

Here, $O(\cdot)$ denotes asymptotic computational complexity. The first term is the filtering cost. The remaining terms are the adaptation and validation costs of the feasible candidates. Filtering before adaptation avoids the update cost of candidates that cannot satisfy the platform-assigned limits.

The final model is $\text{DT}_c^* = (a_{q_c^*}, \theta_{c, q_c^*}^{(T)})$. The platform also records $F_c(a_{q_c^*})$ and $P(a_{q_c^*})$ for the selected architecture.

5. Experimental Evaluation

This section evaluates predictive accuracy, errors on difficult cases, whether selected models satisfy the assigned limits, target-side cost, and generalization. It also examines the main components and settings of RCF-DTI.

Table 1

Main experimental configuration.

Category	Item	Setting
Data	Source / held-out target centers	20 / 20
Data	Observations per center; input length	3000; $L=96$
Data	Main paired cases	80
Target	Horizons; support sizes	H in $\{1, 4\}$; K in $\{10, 20\}$
Target	Validation / check / test windows	80 / 60 / 200
Model	Full architecture space; retained architectures	66 architectures; six retained architectures
Model	Candidate models	7 before filtering
Model	Feasible candidates	4–7 after filtering; mean 6.25
Model	Reference architecture	One-layer GRU; 32 hidden units; dropout 0.1
Model	Reference initialization	Pretrained on pooled source-center data
Adaptation	Optimizer / loss / update budget	SGD / MSE / 50 updates
Adaptation	Learning rate / gradient clipping	0.01 / 1.0
Selection	Selection rule	Select an alternative only when validation WMSE is at least 10% lower than the reference
Selection	Threshold calibration	Disjoint development centers; development check sets only
Model complexity limits	Tight: estimated operation count; parameter count	$1.5e6$ / $3e4$
Model complexity limits	Medium: estimated operation count; parameter count	$5e6$ / $1e5$
Model complexity limits	Loose: estimated operation count; parameter count	$2e7$ / $5e5$
Statistics	Confidence interval	Center-cluster bootstrap; 4000 repeats
Protocol	Test usage	Opened only after the architecture and parameters are fixed

5.1. Experimental Setup

5.1.1. Benchmark and Evaluation Protocol

We use a controlled synthetic benchmark for heterogeneous multi-center workloads. The benchmark varies sampling intervals, missing-data patterns, noise levels, workload dynamics, target data sizes, and model complexity limits. Each center contains 3000 workload observations. The input window length is $L = 96$.

The main evaluation uses 20 source centers and 20 held-out target centers. For each target center, we consider two prediction horizons, $H \in \{1, 4\}$, and two support sizes, $K \in \{10, 20\}$. A target case is one combination of a target center, a prediction horizon, and a support size. The main evaluation therefore contains 80 target cases.

The target centers are grouped into types A, B, and C. These groups differ in sampling intervals, missing-data patterns, noise levels, and workload dynamics. The group label is used only for result analysis. It is not an input to RCF-DTI or any baseline.

We use four disjoint center groups: source, development, diagnostic, and held-out target. Source centers train the source initializations. Development centers are used for architecture retention and threshold calibration. The diagnostic group contains three disjoint pools for the ablation, optimizer-matched, and source-scale studies. Held-out target centers are used only for the main comparison.

Each target case is divided in chronological order into support, validation, check, and test sets. Support data are used for adaptation, and validation data are used for candidate selection. On development cases, validation and check data support architecture retention, and check data also calibrate the threshold. On diagnostic and held-out cases, check data are used only after candidate selection. Test data are used only for final evaluation. Neither check nor test data affect target adaptation or within-case candidate selection.

The reference architecture is a one-layer gated recurrent unit (GRU) with 32 hidden units and a dropout rate of 0.1. Its parameters are pretrained on the pooled source-center data. The resulting source-trained model is used as the reference candidate in all later experiments.

An alternative architecture is retained when it gives lower error than the reference architecture in at least two development cases on either the validation set or the development check set. This rule fixes the compact architecture set before diagnostic and held-out target evaluation.

The selection threshold is fixed on the development cases. A harmful selection occurs when RCF-DTI selects a non-reference candidate whose check loss is higher than that of the reference candidate. We test thresholds of 5%, 7.5%, 10%, 12.5%, 15%, and 20%.

A threshold is eligible only when all development criteria are satisfied. First, its harmful-selection rate must not exceed 5%. Second, its mean check-loss reduction over the reference candidate must be at least 3%, and the lower bound of its 95% confidence interval must be positive. The additional reduction over the reference-architecture control must satisfy the same gain and confidence requirements. The smallest eligible threshold is 10%. The main candidate bank, reference candidate, selection threshold, and other main settings are fixed before the ablation, optimizer-matched, and held-out evaluations. Table 1 lists the main experimental configuration.

5.1.2. Task, Metrics, and Implementation

We use workload prediction only as the evaluation task. It tests whether an instantiated digital twin model captures the temporal behavior of the target center.

Let $\mathbf{X}_{c,t}$ denote an input window of length L . The model predicts the next H workload states as

$$\hat{\mathbf{y}}_{c,t+1:t+H} = f_{a,\theta}(\mathbf{X}_{c,t}), \quad (13)$$

where t is the current time index, $\hat{\mathbf{y}}_{c,t+1:t+H}$ is the predicted workload vector, and $\mathbf{y}_{c,t+1:t+H}$ is the corresponding observed vector.

We use mean squared error (MSE) as the task loss. For a dataset $\mathcal{D}$ of input–output pairs, it is defined as

$$\mathcal{L}(\mathcal{D}; a, \theta) = \frac{1}{|\mathcal{D}|} \sum_{(\mathbf{X}, \mathbf{y}) \in \mathcal{D}} \frac{1}{H} \|f_{a,\theta}(\mathbf{X}) - \mathbf{y}\|_2^2. \quad (14)$$

Here, $|\mathcal{D}|$ is the number of prediction windows, H is the prediction horizon, and $\|\cdot\|_2$ is the Euclidean norm. The loss averages the squared error over all prediction windows and prediction steps. It is the task loss introduced in Section 3.

We report mean absolute error (MAE), MSE, Worst-10% error, conditional value at risk at the 90th percentile (CVaR90), target-side instantiation time, and limit satisfaction. For each target case, MAE and MSE are averaged over all test windows and prediction steps. Worst-10% error is the mean window-level MSE over the 10% of test windows with the largest errors. CVaR90 is the mean case-level MSE over the 10% of target cases with the largest MSE values. Limit satisfaction is the percentage of selected models that meet both assigned model complexity limits.

The main results report averages over the held-out target cases. Since all methods are evaluated on the same target cases, relative gains are first computed for each paired case. The four (H, K) cases from the same target center are treated as one group. The 95% confidence intervals are obtained by resampling target centers 4000 times.

All methods are implemented in Python and PyTorch and run on the same NVIDIA RTX 3060 Laptop GPU. The main source initializations are trained with the Adam optimizer for 50 epochs. The learning rate is 10^{-3} , and the batch size is 64.

Each target update uses the complete support set. The updates are implemented with the SGD optimizer and MSE loss. All feasible RCF-DTI candidates use the same learning rate, gradient clipping rule, and fixed budget of 50 updates. Candidate-specific early stopping is not used.

GPU execution is synchronized before runtime measurement. The timer includes candidate construction, loading of source initializations, target adaptation, validation evaluation, and final candidate selection. It excludes source training, data generation, and post-selection check and test evaluation.

The main candidate bank uses source-training seed 2904. Seeds 2905 and 2906 are used only in the additional training-seed study.

5.1.3. Baselines and Comparison Settings

We compare RCF-DTI with six baselines: pretraining and fine-tuning (PT+FT), MeDeT-based adaptation, random initialization, few-shot neural architecture search, zero-shot neural architecture search, and zero-shot neural architecture search with fine-tuning.

PT+FT adapts the source-trained reference candidate using the same SGD, MSE, and 50-update protocol as RCF-DTI. Random initialization uses the same reference architecture and adaptation protocol without source training. MeDeT-based adaptation follows the meta-initialization idea of MeDeT [11] and uses the reference architecture.

Few-shot neural architecture search follows the candidate-ranking idea in [14]. It retains and adapts the Top-12 candidates. The two zero-shot baselines use architecture scores that do not require full target training [15, 16]. One selects a model without target adaptation. The other adapts its Top-12 candidates before final selection.

The MeDeT and architecture-search baselines are adapted to the common workload prediction setting. They are not full reproductions of the original systems. All methods use the same target cases, chronological splits, complexity limits, and architecture definitions. No method uses the target test set for training, selection, or stopping.

PT+FT is the main controlled baseline because it adapts the same reference candidate as RCF-DTI. We also conduct an optimizer-matched comparison on a separate diagnostic pool. In this comparison, all adapted candidates use SGD,

Table 2

Baseline comparison settings.

Method	Structure	Source initialization	Target update	Adapted models	Complexity-limit check
PT+FT	Fixed reference architecture	Pooled source pretraining	SGD/MSE-50	1	Before output
MeDeT-based adaptation	Fixed reference architecture	Meta initialization	SGD/MSE-50	1	Before output
Random initialization	Fixed reference architecture	Random initialization	SGD/MSE-50	1	Before output
Few-shot NAS	Top-12 shortlist	Baseline-specific initialization bank	Adam/Huber-50	≤ 12	Before candidate evaluation
Zero-shot NAS	Top-ranked model	Baseline-specific initialization	None	0	Before candidate evaluation
Zero-shot NAS + fine-tuning	Top-12 shortlist	Baseline-specific initialization bank	Adam/Huber-50	≤ 12	Before candidate evaluation
RCF-DTI	Source-initialization bank (six architectures)	Architecture-matched source initializations	SGD/MSE-50	4-7	Before adaptation and output

Table 3

Predictive performance on the held-out target centers.

Method	MAE	WMSE	Worst-10%	CVaR90	Limit satisfaction (%)
PT+FT	0.05382	0.005120	0.019078	0.013874	100.0
MeDeT-based adaptation	0.06409	0.007115	0.024170	0.022398	100.0
Random initialization	0.12562	0.026202	0.073298	0.086879	100.0
Few-shot NAS	0.06428	0.007105	0.026501	0.019089	100.0
Zero-shot NAS	0.13500	0.030387	0.075192	0.110634	100.0
Zero-shot NAS + fine-tuning	0.06654	0.007527	0.028539	0.018781	100.0
RCF-DTI	0.04866	0.004219	0.015755	0.011706	100.0

MSE, and 50 updates, and each matched search baseline evaluates 12 candidates. The pool is not used to tune any method. Table 2 summarizes the comparison settings.

5.2. Overall Performance

Table 3 reports the predictive performance on the held-out target cases. All compared methods use the same model complexity limits and the same feasibility check. Therefore, every selected model in this table satisfies both limits. The effect of removing complexity filtering is examined separately in Section 5.5.

RCF-DTI gives the lowest MAE, MSE, Worst-10% error, and CVaR90 among all compared methods. PT+FT is the strongest baseline. Relative to PT+FT, RCF-DTI reduces MAE, MSE, and Worst-10% error by 8.55%, 14.60%, and 13.76%, respectively.

These reductions are first computed for each paired target case and then averaged. They are not computed from the aggregate values in Table 3. The corresponding 95% confidence intervals are [5.42%, 11.85%], [9.41%, 20.14%], and [8.44%, 19.25%]. All three intervals remain above zero.

We next examine the case-level results against PT+FT. Since PT+FT adapts the reference candidate, it also provides the test result of the reference model under the common target adaptation settings. For post-selection analysis, an alternative selection is called beneficial when its test MSE is lower than that of PT+FT. It is called harmful when its test MSE is higher. These labels are assigned after candidate selection using the test set. They do not affect the algorithm. Figure 6 shows the paired comparison. Points below the diagonal favor RCF-DTI.

Table 4

Optimizer-matched comparison on the separate diagnostic target pool.

Method	WMSE	Worst-10%	CVaR90	Adapted models
RCF-DTI	0.003809	0.015094	0.008696	6.10
PT+FT (reference model)	0.004497	0.017144	0.011878	1.00
Few-shot NAS (matched)	0.006715	0.025246	0.019653	12.00
Zero-shot NAS (matched)	0.007402	0.028099	0.020402	12.00
Common 12-candidate shortlist, lowest validation loss	0.006707	0.025240	0.019653	12.00
Common 12-candidate shortlist, minimum-improvement rule	0.006819	0.025562	0.020547	12.00




Figure 6: Paired case-level comparison between RCF-DTI and PT+FT on the held-out target cases. (a) Test MSE. (b) Test Worst-10% error. Points below the diagonal favor RCF-DTI. The beneficial and harmful labels are determined from post-selection test MSE, while reference-retained cases are those in which RCF-DTI keeps the reference candidate. The same outcome categories are used in both panels and do not affect candidate selection.

RCF-DTI selects an alternative in 47 cases. Among them, 44 are beneficial and 3 are harmful. Thus, 93.62% of the alternative selections improve test MSE, while harmful selections account for 3.75% of all held-out cases. The reference is retained in the other 33 cases. Reference regularization therefore reduces, but does not remove, harmful selections.

Table 4 reports the optimizer-matched comparison on the separate diagnostic target pool.

RCF-DTI still gives the lowest MSE, Worst-10% error, and CVaR90. It adapts 6.10 candidates on average, while each matched search baseline adapts 12 candidates. RCF-DTI's advantage therefore remains after matching the optimizer, task loss, and number of target updates.

This comparison does not isolate one component of RCF-DTI. The compared methods still differ in their source initializations and candidate-selection procedures.

5.3. Performance under Target Heterogeneity

We first examine whether the gain of RCF-DTI is limited to one target setting. Figure 7 reports the paired MSE reduction relative to PT+FT under different prediction horizons, support sizes, and target-center types. A positive value means that RCF-DTI gives lower MSE than PT+FT. The confidence intervals are obtained by resampling target centers.

Complexity-Constrained Few-Shot DT Instantiation




Figure 7: Performance across heterogeneous target settings. Points show the mean paired MSE reduction relative to PT+FT. Horizontal bars show 95% confidence intervals obtained by resampling target centers. (a) Results by prediction horizon and support size. (b) Results by target-center type.

RCF-DTI gives positive mean reductions in all four (H, K) settings and all three target-center types. The reductions range from 11.57% to 18.03% across the (H, K) settings and from 13.03% to 16.80% across the center types. All seven confidence intervals remain above zero. Thus, the gain is not limited to one prediction horizon, support size, or center type.

5.4. Effect of Complexity Limits

We next examine how the assigned complexity limits affect target-side instantiation. The benchmark includes tight, medium, and loose limit settings. The tight setting uses the smallest upper bounds on estimated operation count and parameter count. The medium setting uses intermediate bounds, and the loose setting uses the largest bounds.

Figure 8 compares the number of candidates before and after complexity filtering. It also reports the selection rates of the final model configurations under the three limit settings. Here, a model configuration refers to an architecture and its source initialization before target adaptation.




Figure 8: Effect of model complexity limits on target-side instantiation. (a) Mean numbers of candidates before and after complexity filtering. (b) Percentages of held-out cases selecting each final model configuration within the tight, medium, and loose limit settings.

Under the tight setting, filtering reduces the mean candidate count from seven to four. The medium and loose settings keep all seven candidates. Five model configurations are selected across the held-out cases. The tight setting

Table 5

Component ablation on the independent diagnostic target pool.

Variant	WMSE	Worst-10%	Gain vs reference (%)	Harmful alternative selection / all cases (%)	Complexity-feasible outputs (%)
RCF-DTI	0.004063	0.015764	15.79	15.00	100.00
Shared pooled initialization	0.004876	0.019044	1.75	2.50	100.00
Reference candidate only	0.005051	0.019512	0.00	0.00	100.00
Two reference-architecture initializations	0.004917	0.019136	1.29	1.25	100.00
Without minimum-improvement rule	0.004055	0.015690	16.16	17.50	100.00
Without complexity-limit filtering	0.004037	0.015635	16.13	13.75	96.25

selects only GRU configurations, while the medium setting also selects two multilayer perceptron configurations. Under the loose setting, the alternative GRU configurations are selected most often.

These results show that the assigned limits affect both adaptation cost and the final model choice. Because filtering is applied first, every selected model satisfies both limits.

5.5. Ablation Study

Table 5 reports the ablation results on an independent diagnostic target pool. This pool does not affect the main candidate bank, reference candidate, selection threshold, or other main settings. Gain is the paired MSE reduction relative to the reference candidate on the same pool.

We first examine the source-training procedure and architecture choices. The earlier pooled-source-training variant keeps the candidate architectures and the two reference-architecture initializations unchanged. It replaces only the non-reference initializations. This variant, the reference-only variant, and the two-reference-initialization variant all give higher MSE than RCF-DTI. These results indicate that both the source-training procedure and selection across different architectures affect target adaptation. Because the earlier variant changes the complete procedure used to train the non-reference initializations, this comparison does not isolate one source-training factor.

Removing reference regularization slightly reduces mean MSE, but increases the harmful-selection rate from 15.00% to 17.50%. Both rates are higher than the 5% rate used during development calibration. Reference regularization therefore trades a small change in mean MSE for fewer harmful selections, although it does not remove them completely.

Removing complexity filtering slightly reduces mean MSE, but lowers limit satisfaction from 100% to 96.25%. Complexity filtering therefore enforces the assigned hard limits rather than directly reducing prediction error.

5.6. Parameter and Cost Analysis

Figure 9 studies the number of retained architectures, the number of target updates, and the selection threshold. These studies use only development or diagnostic cases.




Figure 9: Parameter studies for RCF-DTI. (a) Diagnostic MSE under different numbers of retained architectures. (b) Check MSE and candidate-selection agreement relative to the 50-update result. (c) Alternative-selection and harmful-selection rates under different selection thresholds. Held-out target cases are not used to choose these settings.

Figure 9(a) shows that diagnostic MSE decreases from 0.004917 with one retained architecture to 0.004063 with four retained architectures, and then remains unchanged. We retain all six architectures because the retained architecture set was fixed before the diagnostic results were observed. This result supports a compact retained architecture set, but does not identify a globally optimal size.

Figure 9(b) shows that check MSE decreases mainly during the first ten updates. Selection agreement with the 50-update result reaches 77.5% at 20 updates. The corresponding rank correlation with the 50-update result is 0.913. We use 50 updates as a fixed common budget, not as a claimed optimal value. Figure 9(c) shows how the alternative-selection and harmful-selection rates change with the threshold. The 10% threshold is the smallest tested value that satisfies all development criteria defined in Section 5.1. It is then fixed for held-out evaluation.

Figure 10 compares MSE, Worst-10% error, CVaR90, target-side time, and model complexity. Each axis is normalized so that the best displayed value receives a score of 100. The raw values remain the basis of the comparison.




Figure 10: Target-side trade-offs of representative methods. The axes show MSE, Worst-10% error, CVaR90, target-side time, parameter count, and estimated operation count. Each axis is normalized across the displayed methods so that the lowest raw value receives a score of 100. Larger normalized scores are better.

RCF-DTI gives the lowest MSE, Worst-10% error, and CVaR90 among the displayed methods, but it is not the fastest. It adapts 6.25 candidates on average and takes 5.676 ± 0.059 seconds per case. Single-model methods are faster because they adapt one candidate. Few-shot NAS and zero-shot NAS with fine-tuning are slower because they adapt 12 candidates. Their target-side times are $2.37\times$ and $3.10\times$ that of RCF-DTI, respectively.

Figure 11 summarizes the final model configurations selected by RCF-DTI. Test results in this figure are used only for post-selection analysis.




Figure 11: Complexity and test performance of the selected model configurations. Estimated operation count is shown on the horizontal axis, paired MSE reduction on the vertical axis, parameter count by marker color, and the number of selected cases by marker area. Circles denote alternative configurations, and the square denotes the reference configuration. The label h shows the number of harmful selections.

The five selected configurations show different operation-count and parameter-count trade-offs. The two multilayer perceptron configurations use fewer estimated operations but more parameters than the GRU configurations. The reference configuration is retained in 33 cases, while four alternative configurations cover the remaining 47 cases. Their mean paired MSE reductions range from 17.43% to 30.50%.

All four alternatives give positive mean reductions, so the overall gain is not produced by one fixed architecture. The three harmful selections occur on the two alternative GRU configurations. Reference-regularized selection therefore reduces, but does not remove, harmful selections.

5.7. Generalization Study

We examine RCF-DTI under changes in source-center scale, source-training seed, and workload data. These studies use the same target-side adaptation and selection procedure as the main evaluation.

Source-center scale. We rebuild the candidate bank using 10, 20, 30, 40, and 50 source centers. Each source-trained model receives 2000 source-training updates at every scale. At each scale, RCF-DTI is compared with a reference candidate trained using the same source-center set and the same source-training budget. The evaluation uses the separate source-scale diagnostic target pool.

Figure 12(a) reports the paired MSE reduction relative to the reference candidate trained at the same source scale. The 20-source-center result belongs to this diagnostic study. It is therefore not directly comparable with the 14.60% main result, which is obtained on the held-out target centers.

The mean reduction is positive at all five source scales, but it does not increase monotonically. Under a fixed source-training budget, adding more source centers distributes the same number of updates across more source workloads. More source centers therefore do not always produce a larger target-side gain.

Source-training seeds and external data. We rebuild the candidate bank using source-training seeds 2904, 2905, and 2906. The three banks use the same source centers, architecture set, source-training budget, and data-generation seed. They are evaluated on the same separate target pool with identical target adaptation seeds and the same target-side protocol. The only changed factor is the source-training seed. This study therefore examines the sensitivity of the candidate bank to source-training randomness.

We also apply the fixed RCF-DTI configuration to machine-usage data from Alibaba Cluster Trace v2018 [31]. The workload observations are real, while the complexity limits are generated by a fixed rule. The candidate bank, adaptation procedure, filtering rule, and selection threshold are not retuned.




Figure 12: Generalization under changed source conditions and Alibaba v2018. (a) Mean paired MSE reduction relative to a reference trained at the same source scale. Shading shows 95% confidence intervals. The vertical dotted line marks the main source scale, and the horizontal dashed line marks zero gain. (b) Case-level reductions for three source-training seeds and Alibaba v2018. The same target cases and target adaptation seeds are used for all three source-training seeds. Violin shapes show distributions, boxes show medians and interquartile ranges, dots show cases, and diamonds with bars show means and 95% confidence intervals. Four Alibaba cases with reductions below -25% are clipped to the left plotting boundary only for visualization; their original values are retained in the released data.

Across the three source-training seeds, the mean paired MSE reduction ranges from 10.52% to 13.21%, and all three confidence intervals remain above zero. Some cases still have zero or negative gains.

On Alibaba v2018, the mean reduction is 1.99%, but its confidence interval includes zero. The distribution is also wider, with four cases below -25% . Thus, the fixed RCF-DTI configuration can be applied without retuning, but it does not show a consistent improvement on this dataset.

6. Conclusion

This paper studied complexity-constrained few-shot digital twin model instantiation for heterogeneous multi-center computing networks. We formulated joint architecture selection and parameter adaptation under limits on inference computation and model size, and proposed RCF-DTI. The algorithm uses a compact source-trained candidate bank, removes infeasible candidates, adapts the remaining candidates with the same target update budget, and applies reference-regularized selection. Experiments show that RCF-DTI reduces MSE by 14.60% relative to pretraining and fine-tuning, while all selected models satisfy both limits. Positive gains remain across heterogeneous target settings and source-training seeds. The fixed configuration can also be applied to Alibaba v2018, although the gain is not statistically clear. Future work will study more operational datasets and measure latency, memory use, and energy consumption on real platforms.

CRedit Authorship Contribution Statement

The final author contribution statement will be added after the author list is confirmed.

Funding

Funding information will be added after confirmation.

Declaration of Competing Interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data Availability

The source code, frozen configurations, processed tables, figure data, and reproduction scripts are available at <https://github.com/w924871681/paper-code-DT-Twin>. The manuscript version is archived in Release v1.1.6: <https://github.com/w924871681/paper-code-DT-Twin/releases/tag/v1.1.6>. The release includes checksums and the files needed to reproduce Tables 1–5 and Figures 1–12. Figures 1–5 are checksum-bound final manuscript assets; Figures 6–12 are regenerated from the released plot-ready data by one public entry point. Synthetic data can be regenerated using the released simulator, configurations, and seeds. The anonymized case-level data used in Figure 12(b) are also included in the release.

The original Alibaba Cluster Trace v2018 is available from Alibaba Group and is not redistributed. The repository provides its source, checksum, file layout, preprocessing procedure, and evaluation entry point. Released Alibaba-derived records contain only anonymized identifiers and processed evaluation values.

References

- [1] L. U. Khan, Z. Han, W. Saad, E. Hossain, M. Guizani, and C. S. Hong, "Digital twin of wireless systems: Overview, taxonomy, challenges, and opportunities," *IEEE Communications Surveys & Tutorials*, vol. 24, no. 4, pp. 2230–2254, Fourth Quarter 2022, doi: 10.1109/COMST.2022.3198273.
- [2] S. Mihai, M. Yaqoob, D. V. Hung, W. Davis, P. Towakel, M. Raza, M. Karamanoglu, B. Barn, D. Shetve, R. V. Prasad, H. Venkataraman, R. Trestian, and H. X. Nguyen, "Digital twins: A survey on enabling technologies, challenges, trends and future prospects," *IEEE Communications Surveys & Tutorials*, vol. 24, no. 4, pp. 2255–2291, Fourth Quarter 2022, doi: 10.1109/COMST.2022.3208773.
- [3] H. Xu, J. Wu, Q. Pan, X. Guan, and M. Guizani, "A survey on digital twin for Industrial Internet of Things: Applications, technologies and tools," *IEEE Communications Surveys & Tutorials*, vol. 25, no. 4, pp. 2569–2598, Fourth Quarter 2023, doi: 10.1109/COMST.2023.3297395.
- [4] F. Tao, H. Zhang, A. Liu, and A. Y. C. Nee, "Digital twin in industry: State-of-the-art," *IEEE Transactions on Industrial Informatics*, vol. 15, no. 4, pp. 2405–2415, 2019, doi: 10.1109/TII.2018.2873186.
- [5] B. Qin, H. Pan, Y. Dai, X. Si, X. Huang, C. Yuen, and Y. Zhang, "Machine and deep learning for digital twin networks: A survey," *IEEE Internet of Things Journal*, vol. 11, no. 21, pp. 34694–34716, Nov. 2024, doi: 10.1109/IIOT.2024.3416733.
- [6] A. Hakiri, A. Gokhale, S. Ben Yahia, and N. Mellouli, "A comprehensive survey on digital twin for future networks and emerging Internet of Things industry," *Computer Networks*, vol. 244, Art. no. 110350, 2024, doi: 10.1016/j.comnet.2024.110350.
- [7] Y. Wu, K. Zhang, and Y. Zhang, "Digital twin networks: A survey," *IEEE Internet of Things Journal*, vol. 8, no. 18, pp. 13789–13804, Sept. 2021, doi: 10.1109/IIOT.2021.3079510.
- [8] J. Li, S. Guo, W. Liang, J. Wang, Q. Chen, Y. Zeng, B. Ye, and X. Jia, "Digital twin-enabled service provisioning in edge computing via continual learning," *IEEE Transactions on Mobile Computing*, vol. 23, no. 6, pp. 7335–7350, 2024, doi: 10.1109/TMC.2023.3332668.
- [9] A. Vettorizzo, M.-R. Bouguelia, J. Vanschoren, T. Rognvaldsson, and K. C. Santosh, "Advances and challenges in meta-learning: A technical review," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 7, pp. 4763–4779, 2024, doi: 10.1109/TPAMI.2024.3357847.
- [10] H. Gharoun, F. Momenifar, F. Chen, and A. H. Gandomi, "Meta-learning approaches for few-shot learning: A survey of recent advances," *ACM Computing Surveys*, vol. 56, no. 12, Art. no. 294, pp. 1–41, 2024, doi: 10.1145/3659943.
- [11] H. Sartaj, S. Ali, and J. M. Gjølby, "MeDeT: Medical device digital twins creation with few-shot meta-learning," *ACM Transactions on Software Engineering and Methodology*, vol. 34, no. 6, Art. no. 158, pp. 1–36, 2025, doi: 10.1145/3708534.
- [12] M. Tan, B. Chen, R. Pang, V. Vasudevan, M. Sandler, A. Howard, and Q. V. Le, "MnasNet: Platform-aware neural architecture search for mobile," in *Proc. IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 2820–2828, doi: 10.1109/CVPR.2019.00293.
- [13] B. Wu, X. Dai, P. Zhang, Y. Wang, F. Sun, Y. Wu, Y. Tian, P. Vajda, Y. Jia, and K. Keutzer, "FBNet: Hardware-aware efficient ConvNet design via differentiable neural architecture search," in *Proc. IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 10734–10742, doi: 10.1109/CVPR.2019.01099.
- [14] Y. Zhao, L. Wang, Y. Tian, R. Fonseca, and T. Guo, "Few-shot neural architecture search," in *Proc. 38th Int. Conf. Mach. Learn. (ICML)*, vol. 139 of *Proc. Mach. Learn. Res.*, pp. 12707–12718, 2021.
- [15] G. Li, D. Hoang, K. Bhardwaj, M. Lin, Z. Wang, and R. Marculescu, "Zero-shot neural architecture search: Challenges, solutions, and opportunities," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 46, no. 12, pp. 7618–7635, Dec. 2024, doi: 10.1109/TPAMI.2024.3395423.

- [16] M.-T. Wu, H.-I. Lin, and C.-W. Tsai, "A training-free neural architecture search algorithm based on search economics," *IEEE Transactions on Evolutionary Computation*, vol. 28, no. 2, pp. 445–459, Apr. 2024, doi: 10.1109/TEVC.2023.3264533.
- [17] Y.-M. Kim, S. Song, B.-M. Koo, J. Son, Y. Lee, and J.-G. Baek, "Enhancing long-term cloud workload forecasting framework: Anomaly handling and ensemble learning in multivariate time series," *IEEE Transactions on Cloud Computing*, vol. 12, no. 2, pp. 789–799, Apr.–Jun. 2024, doi: 10.1109/TCC.2024.3400859.
- [18] F. Zhao, W. Lin, S. Lin, S. Tang, and K. Li, "MSCNet: Multi-scale network with convolutions for long-term cloud workload prediction," *IEEE Transactions on Services Computing*, vol. 18, no. 2, pp. 969–982, Mar.–Apr. 2025, doi: 10.1109/TSC.2025.3536313.
- [19] B. Feng and Z. Ding, "Application-oriented cloud workload prediction: A survey and new perspectives," *Tsinghua Science and Technology*, vol. 30, no. 1, pp. 34–54, Feb. 2025, doi: 10.26599/TST.2024.9010024.
- [20] Y. Tashiro, J. Song, Y. Song, and S. Ermon, "CSDI: Conditional score-based diffusion models for probabilistic time series imputation," in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, pp. 24804–24816, 2021.
- [21] Y. Zhu, W. Yu, and X. Li, "A multi-objective transfer learning framework for time series forecasting with Concept Echo State Networks," *Neural Networks*, vol. 186, Art. no. 107272, 2025, doi: 10.1016/j.neunet.2025.107272.
- [22] L. M. Zintgraf, K. Shiarlis, V. Kurin, K. Hofmann, and S. Whiteson, "Fast context adaptation via meta-learning," in *Proc. 36th Int. Conf. Mach. Learn. (ICML)*, vol. 97 of *Proc. Mach. Learn. Res.*, pp. 7693–7702, 2019.
- [23] Y. Zhao, T. Zhang, J. Li, and Y. Tian, "Dual adaptive representation alignment for cross-domain few-shot learning," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 10, pp. 11720–11732, 2023, doi: 10.1109/TPAMI.2023.3272697.
- [24] J. Wu, X. Liu, X. Yin, T. Zhang, and Y. Zhang, "Task-adaptive prompted transformer for cross-domain few-shot learning," in *Proc. AAAI Conf. Artif. Intell.*, vol. 38, no. 6, pp. 6012–6020, 2024, doi: 10.1609/aaai.v38i6.28416.
- [25] W. Wang, L. Duan, Y. Wang, J. Fan, and Z. Zhang, "MMT: Cross domain few-shot learning via meta-memory transfer," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 45, no. 12, pp. 15018–15035, Dec. 2023, doi: 10.1109/TPAMI.2023.3306352.
- [26] J. Son, S. Lee, and G. Kim, "When meta-learning meets online and continual learning: A survey," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 47, no. 1, pp. 413–432, 2025, doi: 10.1109/TPAMI.2024.3463709.
- [27] Z. Zuo, Y. Huang, Z. Li, Y. Jiang, and C. Liu, "Mixed contrastive transfer learning for few-shot workload prediction in the cloud," *Computing*, vol. 107, no. 1, Art. no. 5, 2025, doi: 10.1007/s00607-024-01366-y.
- [28] Y. Zhao, L. Wang, and T. Guo, "Multi-objective neural architecture search by learning search space partitions," *Journal of Machine Learning Research*, vol. 25, no. 177, pp. 1–41, 2024.
- [29] M. E. Khansari and S. Sharifian, "A deep reinforcement learning approach towards distributed Function as a Service (FaaS) based edge application orchestration in cloud-edge continuum," *Journal of Network and Computer Applications*, vol. 233, Art. no. 104042, 2025, doi: 10.1016/j.jnca.2024.104042.
- [30] D. Hortelano, I. de Miguel, R. J. Durán Barroso, J. C. Aguado, N. Merayo, L. Ruiz, A. Asensio, X. Masip-Bruin, P. Fernández, R. M. Lorenzo, and E. J. Abril, "A comprehensive survey on reinforcement-learning-based computation offloading techniques in Edge Computing Systems," *Journal of Network and Computer Applications*, vol. 216, Art. no. 103669, 2023, doi: 10.1016/j.jnca.2023.103669.
- [31] [dataset] Alibaba Group, "Alibaba Cluster Trace Program: Cluster Trace v2018," GitHub repository, 2018. [Online]. Available: <https://github.com/alibaba/clusterdata/tree/master/cluster-trace-v2018>. Accessed: Jul. 3, 2026.